

Mayank Goel

+91 789-054-5454 | themayankgoel28@gmail.com | [linkedin](#) | [github](#) | [substack](#)

EDUCATION

International Institute of Information Technology - Hyderabad

B.Tech in Computer Science, CGPA: 7.96/10

Hyderabad, India

Jul. 2019 – Apr. 2023

International Institute of Information Technology - Hyderabad

MS in Computational Linguistics, CGPA: 7.92/10

Hyderabad, India

May 2023 – Dec. 2024

EXPERIENCE

Researcher

Lossfunk

Aug. 2025 - Present

Bangalore, India

- Exploring new paradigms to approach AI Safety. Primary topics of interest are moral philosophy, game theory, multi-agent systems.
- Built a transformers model on the Moral Machine Dataset by MIT using a custom architecture. Did mechanistic interpretability of the model to learn more about how humans do moral decision making. This work was accepted at AAAI '26 Workshop on Machine Ethics.
- Explored multiple research questions - token kinematics in LLMs, representation learning, moral embeddings.

Research Mentor

Algoverse Research

Jul. 2025 - Dec. 2025

Remote

- Mentoring teams as part of a research bootcamp, consisting primarily of college students and early-stage engineers. Responsibilities include holding meets, discussing and coming up with the direction, unblocking team members when needed, and acting as liaison to senior mentors and organizers.
- Works from this have been accepted to the SEA Workshop, NeurIPS '25, ICLR '26 Main, MALGAI Workshop, ICLR '26 .
- Executed skills about effective conflict management while maintaining strong epistemic hygiene.

Research Engineer

IBM Research

Mar. 2025 - Jul. 2025

Bangalore, India

- Contributed to IBM's Granite LLM team by improving training data quality via preprocessing, synthetic data generation, and ablation studies. Implemented and evaluated methods from recent literature, gaining deep familiarity with SFT data pipelines and constructing task-specific synthetic datasets grounded in research.
- Led the design and implementation of a scalable data-rewriting pipeline using vLLM, building infrastructure for automated, high-throughput inference. Optimized GPU utilization across nodes by managing concurrency, multi-file workflows, and resource sharing to minimize idle time and maximize throughput.

Researcher

AI Safety Camp

Jan. 2025 - Apr. 2025

Remote

- Part of a team selected by Prof. Masaharu Mizumoto of JAIST on a model to improve the morality of LLMs.
- Involved with discussions on philosophical grounding of research methodology, primarily around Virtue Ethics as a moral framework and approaches for moral dilemma generation.
- Tried SFT/DPO on human-labeled moral dilemmas on models to test performance on several morality related benchmarks. Owned the evaluation pipeline for this project.

Research Engineer

Deccan AI

Sep. 2024 – Mar. 2025

Hyderabad, India

- Involved with benchmarking of models on various tasks, across modalities, such as tasks on image, speech, code. From searching for the appropriate metrics to explaining them to client-facing teams.
- Worked on automated QC for data generated by human annotators. Included fine-tuning models on curated data as part of a company repository.
- Took several interviews for part-time and full-time interviews for ML role. Effectively managed two interns to help their skill development and contribution to the company.

Research Intern

ComplexData Lab, MILA AI

Oct. 2023 – Jul. 2024

Remote

- Designed experiments using the Debate framework for uncertainty quantification of LLM outputs.

- Worked on experiments to analyze and visualize data from other team projects on uncertainty quantification.
- Work done during this phase was accepted to COLM '24 and NeurIPS SafeGenAI Workshop '24

Teaching Assistant

Jul. 2021 – Apr. 2024

IIT Hyderabad

Hyderabad, India

- Was involved with designing and evaluating assignments and conducting tutorials for the course Introduction to NLP in Spring '22 and Advanced NLP in Monsoon '23

Research Intern

Jun. 2023 – Jul. 2023

Robert Bosch GmbH

Bangalore, India

- Built upon an existing document hierarchical clustering project, add functionality to use sentence transformers.
- Built a robust pipeline to fine-tune sentence transformers using various methods.
- Developed a novel evaluation metric exclusively tailored to hierarchical clustering, addressing existing limitations with evaluation methods.

Research Associate

Mar. 2021 – Feb. 2023

Proxzar AI

Remote

- Fine-tuned Llama-7B on finance data using LORA and quantization. Did a hyperparameter search.
- Developed a robust e-commerce website crawling pipeline using Apache Nutch and Selenium
- Built a Spacy3 model utilizing Transformers for a production-grade custom NER model for an e-commerce company, as part of an end-to-end system, with a dataset augmented by 20% using fuzzy matching. Achieved accuracy of 91% on a multi-class system

Research Intern

May 2021 – Jun. 2021

Trivedi Centre for Political Data, Ashoka University

Remote

- Built a streamlit dashboard to display visualized data of MPs and Governors in India, which accepts a variety of parameters. Work can be viewed at [\[link\]](#).

PUBLICATIONS

Automated Stateful Specialization for Adaptive Agent Systems ICLR '26

Authors: Myan Vu, Harrish Ayyanar, PANG JIANG, Anwiketh Reddy, **Mayank Goel**

Building Interpretable Models for Moral Decision-Making AAI '26, Machine Ethics Workshop

Authors: **Mayank Goel**, Aritra Das, Paras Chopra

Automated Specialization of Stateful Agent Systems NeurIPS '25 Scaling Environments for Agents Workshop

Authors: Myan Vu, Harrish Ayyanar, Pang Jiang, Anwiketh Reddy, **Mayank Goel**, Kevin Zhu

Conceptual Limitations of Current AI Safety Approaches and Virtue Ethics as an Alternative Technical AI Safety Conference '25

Authors: Masaharu Mizumoto, Rujuta Karekar, Mads Udengaard, **Mayank Goel**, Daan Henselmans, Nurshafira Noh, Saptadip Saha, Pranshu Bohra

Epistemic Integrity in Large Language Models NeurIPS '24 SafeGenAI Workshop

Authors: Bijean Ghafouri, Shahrard Mohammadzadeh, James Zhou, Pratheeksha Nair, Jacob-Junqi Tian, **Mayank Goel**, Reihaneh Rabbany, Jean-François Godbout, Kellin Pelrine

Web Retrieval Agents for Evidence-Based Misinformation Detection COLM '24

Authors: Jacob-Junqi Tian, Hao Yu, Yury Orlovskiy, Tyler Vergho, Mauricio Rivera, **Mayank Goel**, Zachary Yang, Jean-François Godbout, Reihaneh Rabbany, Kellin Pelrine

Joke Generation Using Masked Language Infilling on Automatically Extracted Templates Humor Research Conference '23

Authors: **Mayank Goel**, Abhijit Manatkar

Towards conversational humor analysis and design Humor Research Conference '21

Authors: Tanishq Chaudhary*, **Mayank Goel***, Radhika Mamidi

Automated Evalauation of Conversational Prompts using Ordinal Ranking for Humor Generation ICON '24

Authors: **Mayank Goel**, Parameswari Krishnamurthy, Radhika Mamidi

TECHNICAL SKILLS

Languages: Python (proficient), Bash, Javascript, C++

Frameworks: PyTorch, HuggingFace Transformers, Streamlit, Selenium, SpaCy

AWARDS AND OTHER RELEVANT INFORMATION

- **Research Award** from IIIT Hyderabad in 2021
- **Top 30** in the **Indian** edition of **International Olympiad of Linguistics**
- **Participant**, **ARENA 2024**, Spring
- **Participant**, **AI Safety Fundamentals Course**, BlueDot Impact
- **Ex-head** of **University Debate Society** and **Humor Club**
- **Founder-Moderator** of **Philosophy Friday Discussion Group**
- **Events Head** of **Felicity**, IIIT's student fest.
- **Best Project**, **Theme Track** in **Natural Language API Hackathon** by expert.ai
- **1st Prize**, **Equity Track** in **HackUMBC**, University of Maryland
- **Best Web Application** in **HackAtHome**, Brown University
- **3rd Prize**, **Walmart Track** in **TAMU Datathon**, Texas A&M University
- **1st Prize**, **Stellantis Track** in **Megathon '22**, IIIT Hyderabad